

February 22, 2025

1 Backpack Price Prediction

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

1.1 Reading the dataset

```
[2]: df = pd.read_csv("./train.csv")
```

1.2 Exploring the dataset

```
[3]: print("Information of the dataset: ")
print(df.info())
```

Information of the dataset:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 300000 entries, 0 to 299999

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	id	300000 non-null	int64
1	Brand	290295 non-null	object
2	Material	291653 non-null	object
3	Size	293405 non-null	object
4	Compartments	300000 non-null	float64
5	Laptop Compartment	292556 non-null	object
6	Waterproof	292950 non-null	object
7	Style	292030 non-null	object
8	Color	290050 non-null	object
9	Weight Capacity (kg)	299862 non-null	float64
10	Price	300000 non-null	float64

dtypes: float64(3), int64(1), object(7)

memory usage: 25.2+ MB

None

```
[4]: print("Description of the dataset: ")
      print(df.describe())
```

Description of the dataset:

	id	Compartments	Weight	Capacity (kg)	Price
count	300000.000000	300000.000000		299862.000000	300000.000000
mean	149999.500000	5.443590		18.029994	81.411107
std	86602.684716	2.890766		6.966914	39.039340
min	0.000000	1.000000		5.000000	15.000000
25%	74999.750000	3.000000		12.097867	47.384620
50%	149999.500000	5.000000		18.068614	80.956120
75%	224999.250000	8.000000		24.002375	115.018160
max	299999.000000	10.000000		30.000000	150.000000

```
[5]: print("Shape of the dataset ", df.shape)
```

Shape of the dataset (300000, 11)

1.3 Exploring the null values

```
[6]: df.isnull().sum()
```

```
[6]: id          0
      Brand      9705
      Material   8347
      Size       6595
      Compartments 0
      Laptop Compartment 7444
      Waterproof  7050
      Style      7970
      Color      9950
      Weight Capacity (kg) 138
      Price      0
      dtype: int64
```

1.4 Handling the null values

```
[7]: df['Brand'] = df['Brand'].fillna("Non Branded")

      df['Material'] = df['Material'].fillna("Unknown")

      df['Size'] = df['Size'].fillna('Unknown')

      df['Laptop Compartment'] = df['Laptop Compartment'].fillna("Unknown")

      df['Waterproof'] = df['Waterproof'].fillna("Unknown")
```

```

df['Style'] = df['Style'].fillna("Unknown")

df['Color'] = df['Color'].fillna('Unique')

weight_mean = df['Weight Capacity (kg)'].mean()
df['Weight Capacity (kg)'] = df['Weight Capacity (kg)'].fillna(weight_mean)

print("After removing all the outliers")
df.isnull().sum()

```

After removing all the outliers

```

[7]: id          0
     Brand        0
     Material     0
     Size         0
     Compartments 0
     Laptop Compartment 0
     Waterproof   0
     Style        0
     Color        0
     Weight Capacity (kg) 0
     Price        0
     dtype: int64

```

1.5 Adding Columns for better Graph plotting

```

[8]: bins = [0, 13, 21, 30]
     labels = ['Light', 'Medium', 'Heavy']
     df['Weight Category'] = pd.cut(df['Weight Capacity (kg)'], bins = bins, labels_
     ↪= labels, right = True)

     category_counts = df['Weight Category'].value_counts().reset_index()
     category_counts.columns = ['Weight Category', 'Count']

```

```

[9]: bins = [15, 60, 105, 150]
     labels = ['Affordable', 'Medium', 'Expensive']
     df['Price Category'] = pd.cut(df['Price'], bins = bins, labels = labels, right_
     ↪= True)

```

1.6 Skewness

```

[10]: numerical_cols = df.select_dtypes(include=['float64', 'int64'])
     skewness = numerical_cols.skew()

     # Print skewness values
     print("Skewness of numerical columns:")

```

```

print(skewness)

# Plotting skewness using subplots for numerical columns
n_cols = len(numerical_cols.columns)

# Set up the subplots
fig, axes = plt.subplots(nrows=1, ncols=n_cols, figsize=(5 * n_cols, 4))


# Plot each numerical column
for ax, col in zip(axes, numerical_cols.columns):
    sns.histplot(df[col], ax=ax, kde=True)
    ax.set_title(f'Distribution of {col}\nSkewness: {skewness[col]:.2f}')
    ax.axvline(df[col].mean(), color='red', linestyle='--', label='Mean')
    ax.axvline(df[col].median(), color='blue', linestyle='--', label='Median')
    ax.legend()

plt.tight_layout()
plt.show()

```

Skewness of numerical columns:

id	0.000000
Compartments	0.029125
Weight Capacity (kg)	-0.064254
Price	0.036883
dtype:	float64



```
[11]: df.isnull().sum()
```

```

[11]: id          0
      Brand        0
      Material     0
      Size         0
      Compartments 0
      Laptop Compartment 0
      Waterproof    0
      Style         0
      Color         0
      Weight Capacity (kg) 0

```

```
Price          0
Weight Category 0
Price Category 693
dtype: int64
```

1.7 Univariate Analysis

```
[12]: # Brand
plt.figure(figsize = (8, 5))
sns.countplot(x = df['Brand'], data = df, palette = 'pastel', hue = None)

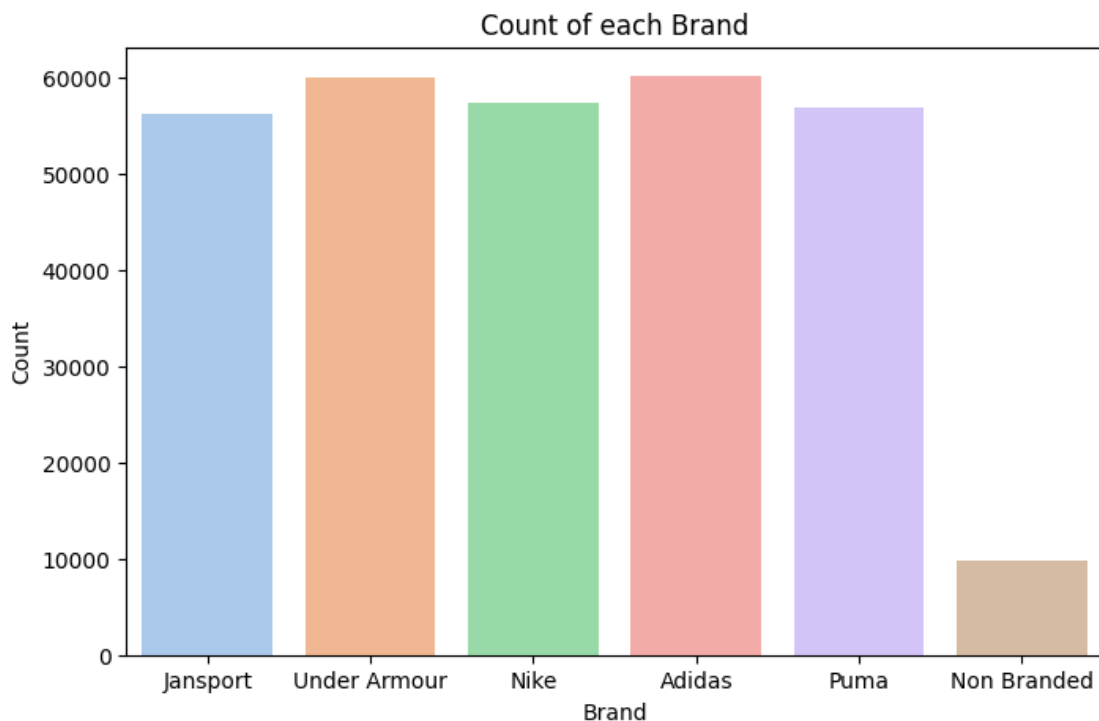
plt.title('Count of each Brand')
plt.xlabel('Brand')
plt.ylabel('Count')

plt.show()
```

<ipython-input-12-4aeb48e43e68>:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x = df['Brand'], data = df, palette = 'pastel', hue = None)
```



```
[13]: # Brand
plt.figure(figsize = (8, 5))
sns.countplot(x = df['Material'], data = df, palette = 'muted', hue = None)


plt.title('Count of each Material')
plt.xlabel('Material')
plt.ylabel('Count')

plt.show()
```

<ipython-input-13-71edb78387b0>:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x = df['Material'], data = df, palette = 'muted', hue = None)
```



```
[14]: # Brand
plt.figure(figsize = (8, 5))
sns.countplot(x = df['Style'], data = df, palette = 'Set1', hue = None)
```

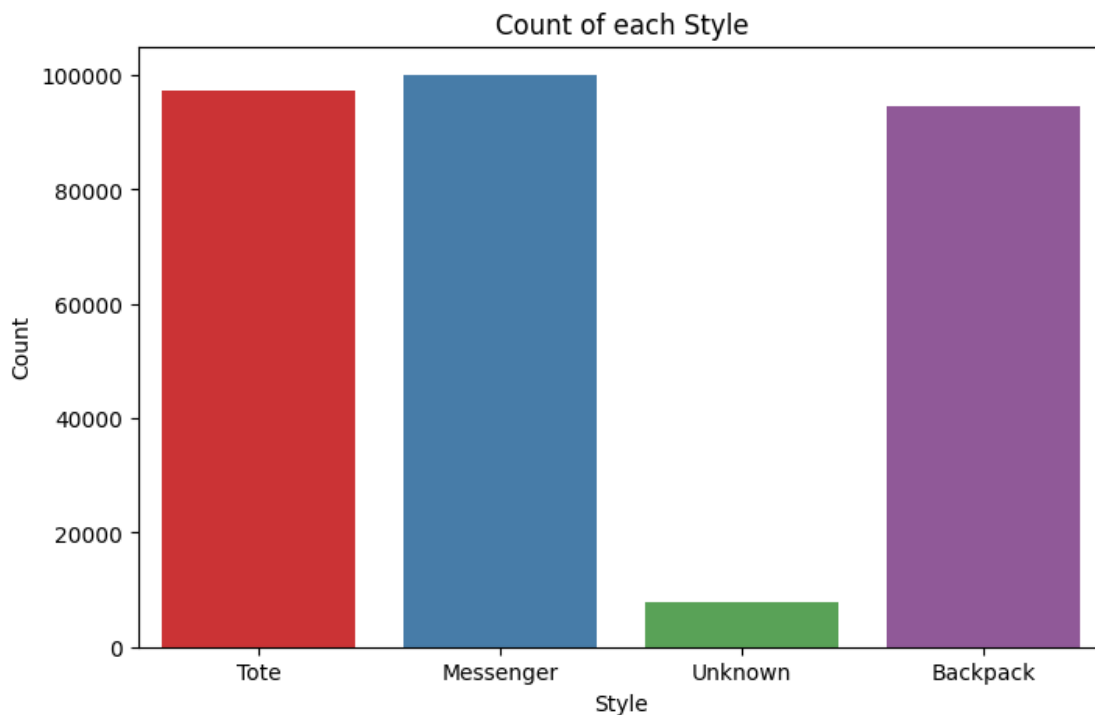
```
plt.title('Count of each Style')
plt.xlabel('Style')
plt.ylabel('Count')

plt.show()
```

<ipython-input-14-eef356250b69>:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x = df['Style'], data = df, palette = 'Set1', hue = None)
```




```
[15]: # Color
color = df.groupby('Color', as_index = False).count()

plt.figure(figsize=(8, 8))
plt.pie(color['Brand'], labels=color['Color'], autopct='%1.1f%%', startangle=90)

# Draw a white circle in the center to create a doughnut effect
centre_circle = plt.Circle((0, 0), 0.70, fc='white')
fig = plt.gcf()
fig.gca().add_artist(centre_circle)
```

```
# Equal aspect ratio ensures that pie is drawn as a circle
plt.axis('equal')
plt.title('Color Distribution')
plt.show()
```



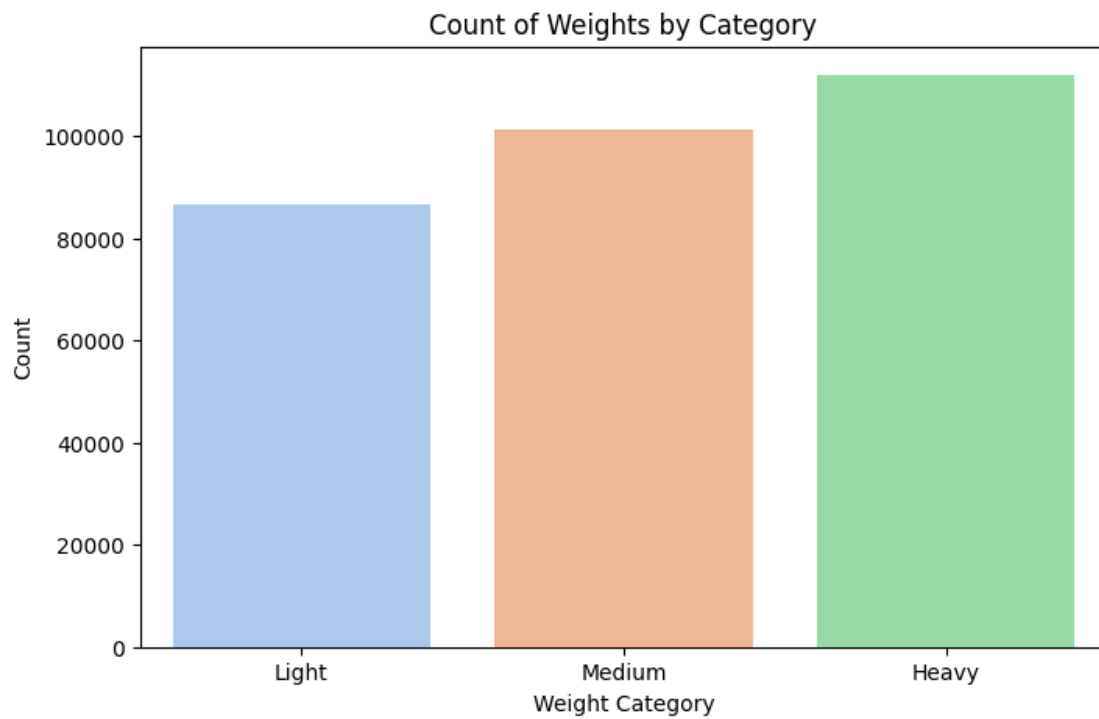
```
[16]: # Weight category
plt.figure(figsize=(8, 5))
sns.barplot(x='Weight Category', y='Count', data=category_counts,
            palette='pastel')
plt.title('Count of Weights by Category')
plt.xlabel('Weight Category')
plt.ylabel('Count')
plt.show()
```



```
<ipython-input-16-2c15eedf876d>:3: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

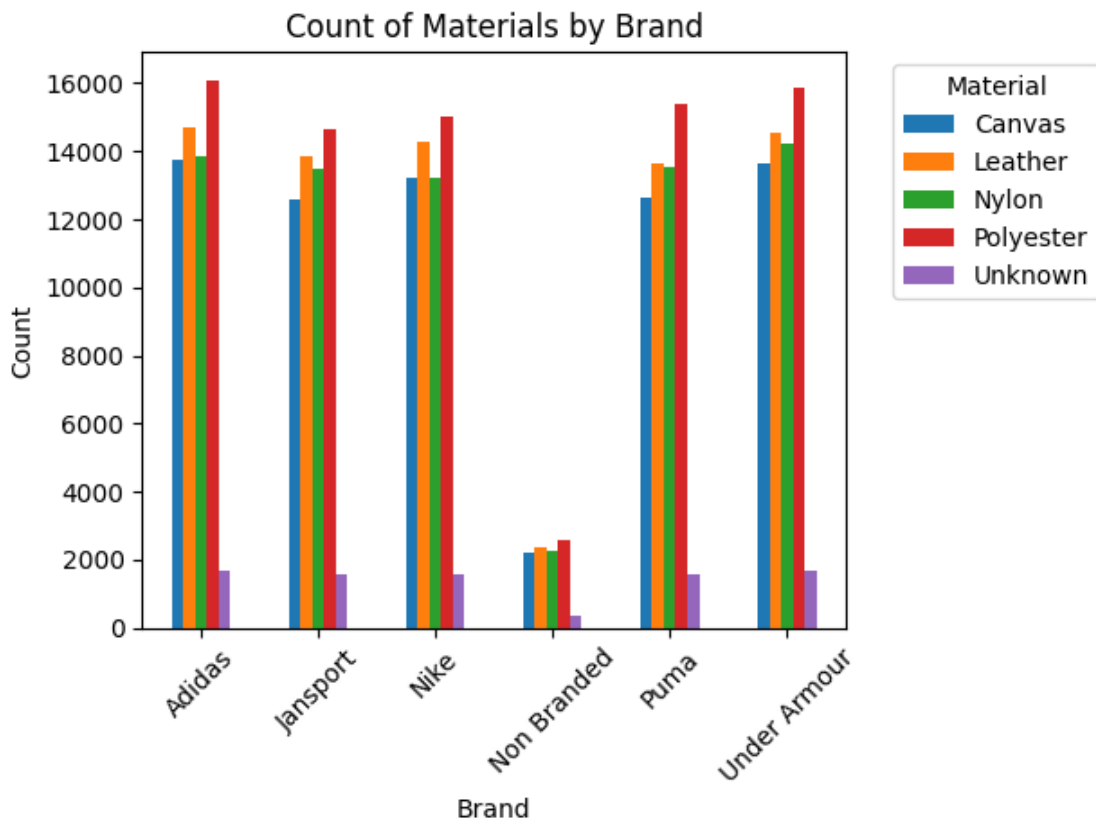
```
sns.barplot(x='Weight Category', y='Count', data=category_counts,  
palette='pastel')
```



1.8 Bivariate analysis

```
[17]: # Brand vs Material
grouped = df.groupby(['Brand', 'Material']).size().reset_index(name='Count')
pivoted = grouped.pivot(index='Brand', columns='Material', values='Count').
    ↪ fillna(0)

pivoted.plot(kind='bar', stacked=False)
plt.title('Count of Materials by Brand')
plt.xlabel('Brand')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.legend(title='Material', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()
```



```
[18]: # Brand vs Price
grouped = df.groupby(['Brand', 'Price Category']).size().
    ↪ reset_index(name='Count')
pivoted = grouped.pivot(index='Brand', columns='Price Category',
    ↪ values='Count').fillna(0)
```

```

pivoted.plot(kind='bar', stacked=False)
plt.title('Count of Brand by Price Category')
plt.xlabel('Brand')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.legend(title='Price Category', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()

```

<ipython-input-18-facb2f982896>:2: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```

grouped = df.groupby(['Brand', 'Price
Category']).size().reset_index(name='Count')

```



```

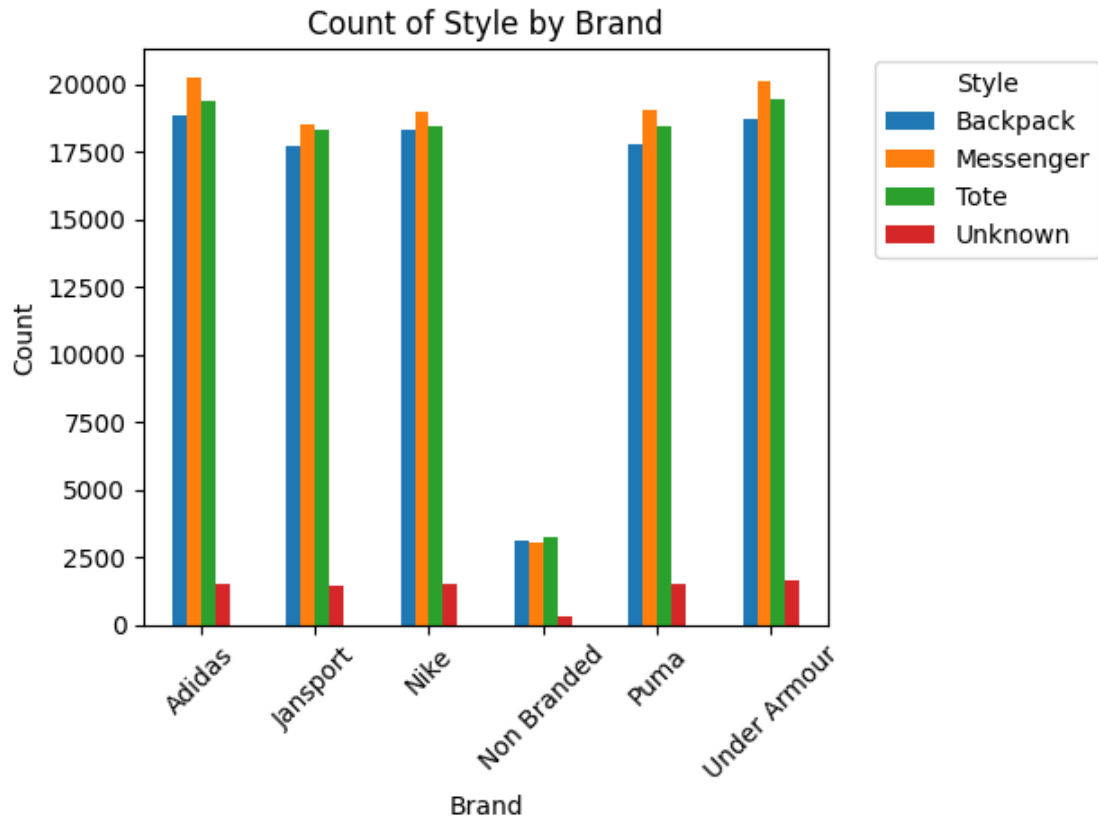
[19]: # Brand vs Style
grouped = df.groupby(['Brand', 'Style']).size().reset_index(name='Count')
pivoted = grouped.pivot(index='Brand', columns='Style', values='Count').
    fillna(0)

```

```

pivoted.plot(kind='bar', stacked=False)
plt.title('Count of Style by Brand')
plt.xlabel('Brand')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.legend(title='Style', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()

```



```

[20]: # Brand vs Weight Category
grouped = df.groupby(['Brand', 'Weight Category']).size().
    ↪reset_index(name='Count')
pivoted = grouped.pivot(index='Brand', columns='Weight Category',
    ↪values='Count').fillna(0)

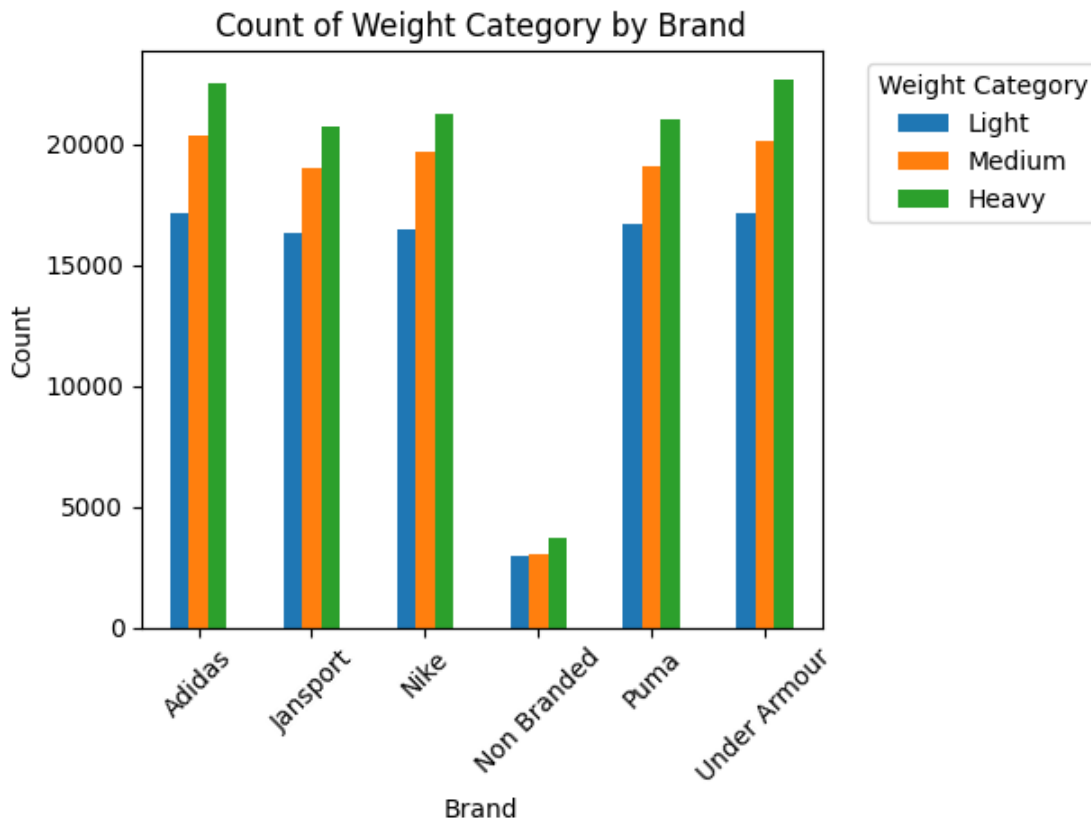
pivoted.plot(kind='bar', stacked=False)
plt.title('Count of Weight Category by Brand')
plt.xlabel('Brand')
plt.ylabel('Count')
plt.xticks(rotation=45)

```

```
plt.legend(title='Weight Category', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()
```

<ipython-input-20-ca41bccade79>:2: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
grouped = df.groupby(['Brand', 'Weight
Category']).size().reset_index(name='Count')
```



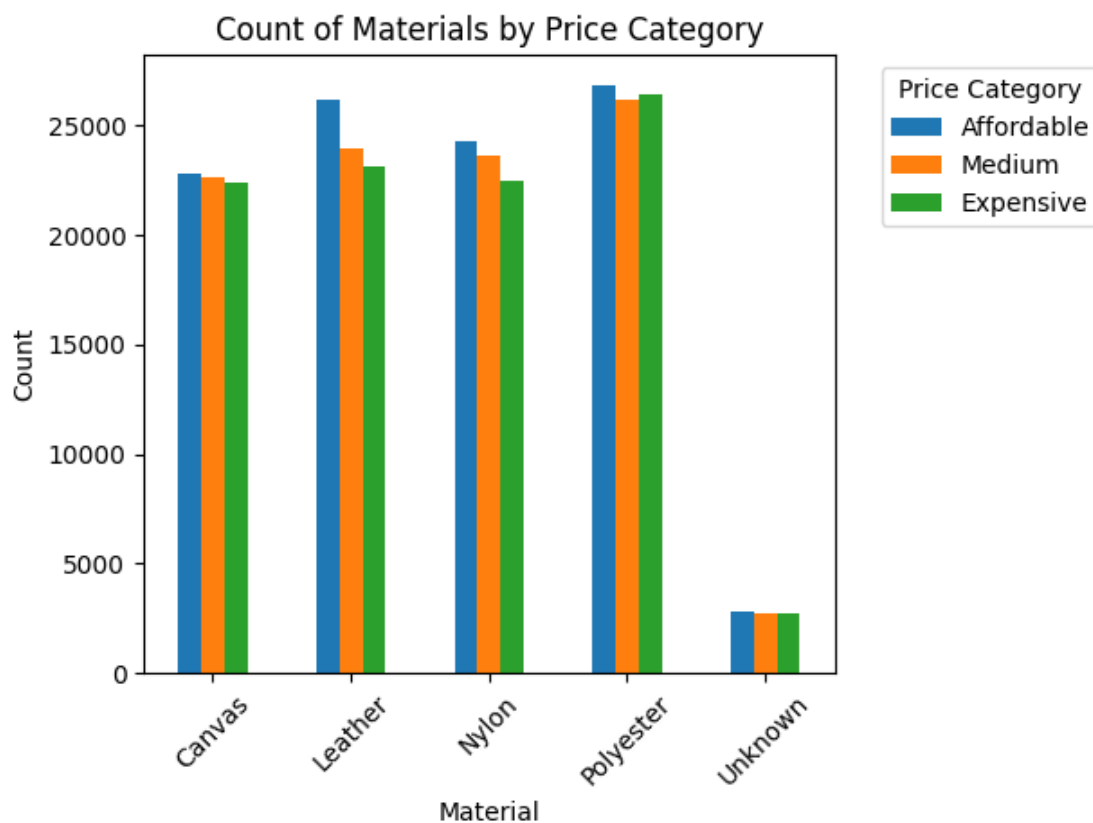
```
[21]: # Material vs Price
grouped = df.groupby(['Material', 'Price Category']).size().
    ↪reset_index(name='Count')
pivoted = grouped.pivot(index='Material', columns='Price Category',
    ↪values='Count').fillna(0)

pivoted.plot(kind='bar', stacked=False)
plt.title('Count of Materials by Price Category')
plt.xlabel('Material')
```

```
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.legend(title='Price Category', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()
```

<ipython-input-21-aa9cd85781be>:2: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
grouped = df.groupby(['Material', 'Price
Category']).size().reset_index(name='Count')
```



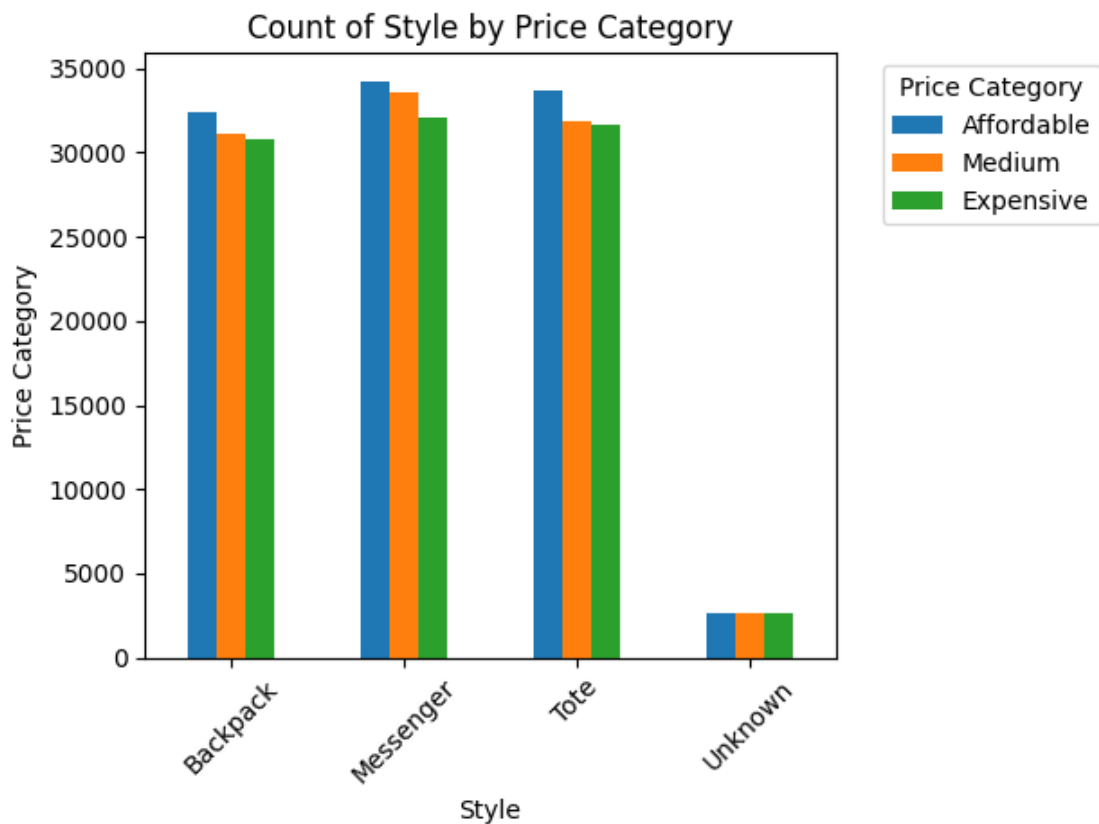
```
[22]: # Style vs Price
grouped = df.groupby(['Style', 'Price Category']).size().
    ↪reset_index(name='Count')
pivoted = grouped.pivot(index='Style', columns='Price Category',
    ↪values='Count').fillna(0)

pivoted.plot(kind='bar', stacked=False)
```

```
plt.title('Count of Style by Price Category')
plt.xlabel('Style')
plt.ylabel('Price Category')
plt.xticks(rotation=45)
plt.legend(title='Price Category', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()
```

<ipython-input-22-421eda627aae>:2: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
grouped = df.groupby(['Style', 'Price
Category']).size().reset_index(name='Count')
```



```
[23]: # Color vs Price
grouped = df.groupby(['Color', 'Price Category']).size().
    ↪reset_index(name='Count')
pivoted = grouped.pivot(index='Color', columns='Price Category',
    ↪values='Count').fillna(0)
```

```

pivoted.plot(kind='bar', stacked=False)
plt.title('Count of Color by Price Category')
plt.xlabel('Color')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.legend(title='Price Category', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()

```

<ipython-input-23-555c3cb96ee6>:2: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```

grouped = df.groupby(['Color', 'Price
Category']).size().reset_index(name='Count')

```



1.9 Feature Engineering for Model Training

```
[24]: # Drop unwanted columns
new_df = df.drop(['id', 'Waterproof', 'Laptop Compartment', 'Weight Category',
↪ 'Price Category'], axis = 1)
```

```
[25]: # Change the float values to int as some models give error for float values
new_df['Compartments'] = new_df['Compartments'].astype(int)
new_df['Price'] = new_df['Price'].astype(int)
new_df['Weight Capacity (kg)'] = new_df['Weight Capacity (kg)'].astype(int)
```

1.10 Model Training

```
[26]: # Defining X and Y dataset
X = new_df.drop(columns = ['Price'])
Y = new_df['Price']
```

```
[27]: # Data preprocessing
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,Y,test_size=0.
↪ 15,random_state=2)
```

1.10.1 Import libraries

```
[28]: from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
from sklearn.metrics import r2_score,mean_absolute_error

from xgboost import XGBRegressor
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.cluster import KMeans, DBSCAN, AgglomerativeClustering
from sklearn.ensemble import
↪ RandomForestRegressor, GradientBoostingRegressor, AdaBoostRegressor, ExtraTreesRegressor
from sklearn.ensemble import VotingRegressor, BaggingRegressor,
↪ StackingRegressor
from sklearn.svm import SVR
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
from sklearn.model_selection import train_test_split
```

1.11 Linear Regression

```
[29]: step1 = ColumnTransformer(  
    ↳ transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0,1,2,4,5])],  
    remainder='passthrough'  
)  
  
step2 = LinearRegression()  
  
pipe = Pipeline([  
    ('step1', step1),  
    ('step2', step2)  
)  
  
pipe.fit(X_train, y_train)  
  
y_pred = pipe.predict(X_test)  
  
print('RMSE score', np.sqrt(mean_squared_error(y_test, y_pred)))
```

RMSE score 39.18521675416943

1.12 Ridge Regression

```
[30]: step1 = ColumnTransformer(  
    ↳ transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0,1,2,4,5])],  
    remainder='passthrough'  
)  
  
step2 = Ridge()  
  
pipe = Pipeline([  
    ('step1', step1),  
    ('step2', step2)  
)  
  
pipe.fit(X_train, y_train)  
  
y_pred = pipe.predict(X_test)  
  
print('RMSE score', np.sqrt(mean_squared_error(y_test, y_pred)))
```

RMSE score 39.18521621559453

1.13 Lasso Regression

```
[31]: step1 = ColumnTransformer(  
    ↳  
    ↳transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0, 1, 2, 4, 5])],  
    ↳remainder='passthrough'  
    )  
  
    step2 = Lasso()  
  
    pipe = Pipeline([  
        ('step1', step1),  
        ('step2', step2)  
    ])  
  
    pipe.fit(X_train, y_train)  
  
    y_pred = pipe.predict(X_test)  
  
    print('RMSE score', np.sqrt(mean_squared_error(y_test, y_pred)))
```

RMSE score 39.19518883212979

1.14 ElasticNet Regression

```
[32]: step1 = ColumnTransformer(  
    ↳  
    ↳transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0, 1, 2, 4, 5])],  
    ↳remainder='passthrough'  
    )  
  
    step2 = ElasticNet()  
  
    pipe = Pipeline([  
        ('step1', step1),  
        ('step2', step2)  
    ])  
  
    pipe.fit(X_train, y_train)  
  
    y_pred = pipe.predict(X_test)  
  
    print('RMSE score', np.sqrt(mean_squared_error(y_test, y_pred)))
```

RMSE score 39.19413509444018

1.15 Decision Tree

```
[33]: step1 = ColumnTransformer(
    ↳
    ↳ transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0,1,2,4,5])],
    ↳ remainder='passthrough'
    )

step2 = DecisionTreeRegressor()

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('RMSE score', np.sqrt(mean_squared_error(y_test, y_pred)))
```

RMSE score 53.4981778839449

1.16 Extra tree

```
[34]: step1 = ColumnTransformer(
    ↳
    ↳ transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0,1,2,4,5])],
    ↳ remainder='passthrough'
    )

step2 = ExtraTreesRegressor(n_estimators=100,
                             random_state=3,
                             max_samples=None,
                             max_features=0.75,
                             max_depth=15)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('RMSE score', np.sqrt(mean_squared_error(y_test, y_pred)))
```

RMSE score 39.49093458977512

1.17 Voting

```
[35]: step1 = ColumnTransformer(
        transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'),
        ↪[0, 1, 2, 4, 5])],
        remainder='passthrough'
    )

    ridge_model = Ridge()
    linear_model = LinearRegression()
    lasso_model = Lasso()

    voting_regressor = VotingRegressor(estimators=[
        ('ridge', ridge_model),
        ('linear', linear_model),
        ('lasso', lasso_model)
    ])

    pipe = Pipeline([
        ('step1', step1),
        ('step2', voting_regressor)
    ])

    pipe.fit(X_train, y_train)

    y_pred = pipe.predict(X_test)

    print('RMSE:', np.sqrt(mean_squared_error(y_test, y_pred)))
```

RMSE: 39.184375647684455

1.18 Bagging

```
[36]: step1 = ColumnTransformer(
        transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'),
        ↪[0, 1, 2, 4, 5])],
        remainder='passthrough'
    )

    base_model = Ridge() # You can change this to LinearRegression() or
    ↪DecisionTreeRegressor()

    bagging_regressor = BaggingRegressor(estimator=base_model, n_estimators=50,
    ↪random_state=42)
```

```

pipe = Pipeline([
    ('step1', step1),
    ('step2', bagging_regressor)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('RMSE:', np.sqrt(mean_squared_error(y_test, y_pred)))

```

RMSE: 39.184898304234295

1.19 Random Forest

```

[37]: step1 = ColumnTransformer(
    □
    ↪transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0,1,2,4,5])],
    remainder='passthrough'
)

step2 = RandomForestRegressor(n_estimators = 100, random_state = 42)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('RMSE', np.sqrt(mean_squared_error(y_test, y_pred)))

```

RMSE 43.9357107695224

1.20 AdaBoost

```

[38]: step1 = ColumnTransformer(
    □
    ↪transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0,1,2,4,5,6])],
    remainder='passthrough'
)

step2 = AdaBoostRegressor(n_estimators=50, random_state=42)

pipe = Pipeline([

```

```

        ('step1',step1),
        ('step2',step2)
    ])

    pipe.fit(X_train,y_train)

    y_pred = pipe.predict(X_test)

    print('RMSE', np.sqrt(mean_squared_error(y_test, y_pred)))

```

RMSE 39.207410208141226

1.21 K Means Clustering

```

[39]: step1 = ColumnTransformer(
        □
        ↪transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0,1,2,4,5,6])],
        remainder='passthrough'
    )

    step2 = KMeans(n_clusters = 3, random_state = 42)

    pipe = Pipeline([
        ('step1',step1),
        ('step2',step2)
    ])

    pipe.fit(X_train,y_train)

    y_pred = pipe.predict(X_test)

    print('RMSE', np.sqrt(mean_squared_error(y_test, y_pred)))

```

RMSE 89.14747095559008

1.22 Gradient Boosting

```

[40]: step1 = ColumnTransformer(
        transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), □
        ↪[0, 1, 2, 4, 5, 6])],
        remainder='passthrough'
    )

    step2 = GradientBoostingRegressor(random_state=42)

    pipe = Pipeline([
        ('step1', step1),

```

```

        ('step2', step2)
    ])

    pipe.fit(X_train, y_train)

    y_pred = pipe.predict(X_test)

    print('RMSE', np.sqrt(mean_squared_error(y_test, y_pred)))

```

RMSE 39.187112181786034

1.23 Stacking

```

[41]: step1 = ColumnTransformer(
        transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'),
        ↪[0, 1, 2, 4, 5, 6])],
        remainder='passthrough'
    )

    base_learners = [
        ('lr', LinearRegression()),
        ('gb', GradientBoostingRegressor(random_state=42)),
    ]

    final_estimator = GradientBoostingRegressor(random_state=42)

    stacking_model = StackingRegressor(estimators=base_learners,
    ↪final_estimator=final_estimator)

    pipe = Pipeline([
        ('step1', step1),
        ('stacking', stacking_model)
    ])

    pipe.fit(X_train, y_train)

    y_pred = pipe.predict(X_test)

    print('RMSE', np.sqrt(mean_squared_error(y_test, y_pred)))

```

RMSE 39.18726052784003

1.24 K Nearest Neighbors

```
[42]: step1 = ColumnTransformer(  
    ↳  
    ↳transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0, 1, 2, 4, 5, 6])],  
    remainder='passthrough'  
)  
  
step2 = KNeighborsRegressor(n_neighbors = 5)  
  
pipe = Pipeline([  
    ('step1', step1),  
    ('step2', step2)  
])  
  
pipe.fit(X_train, y_train)  
  
y_pred = pipe.predict(X_test)  
  
print('RMSE', np.sqrt(mean_squared_error(y_test, y_pred)))
```

RMSE 42.959242254852576

1.25 XG Boost

```
[43]: step1 = ColumnTransformer(  
    ↳  
    ↳transformers=[('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0, 1, 2, 4, 5, 6])],  
    remainder='passthrough'  
)  
  
step2 = XGBRegressor(random_state=42)  
  
pipe = Pipeline([  
    ('step1', step1),  
    ('step2', step2)  
])  
  
pipe.fit(X_train, y_train)  
  
y_pred = pipe.predict(X_test)  
  
print('RMSE', np.sqrt(mean_squared_error(y_test, y_pred)))
```

RMSE 39.33804068980177